

# amazon

**MMA865 Final Project**  
**Team Kipling**



# 01 Problem & Objective

# The Business Problem and Objective

- **Business Problem:** According to research<sup>1</sup>, 91 percent of people regularly or occasionally read online reviews, and 84 percent trust online reviews as much as personal recommendations. Amazon's customers rely on them to ensure that they're getting what they're paying for. However, not all reviews are legitimate or add value to their decision-making process.
- **Business Objective:** How can Amazon accurately identify and leverage helpful reviews to improve customer experience and maintain customer loyalty?

1- Source: Brightlocal Local Consumer Review Survey 2020

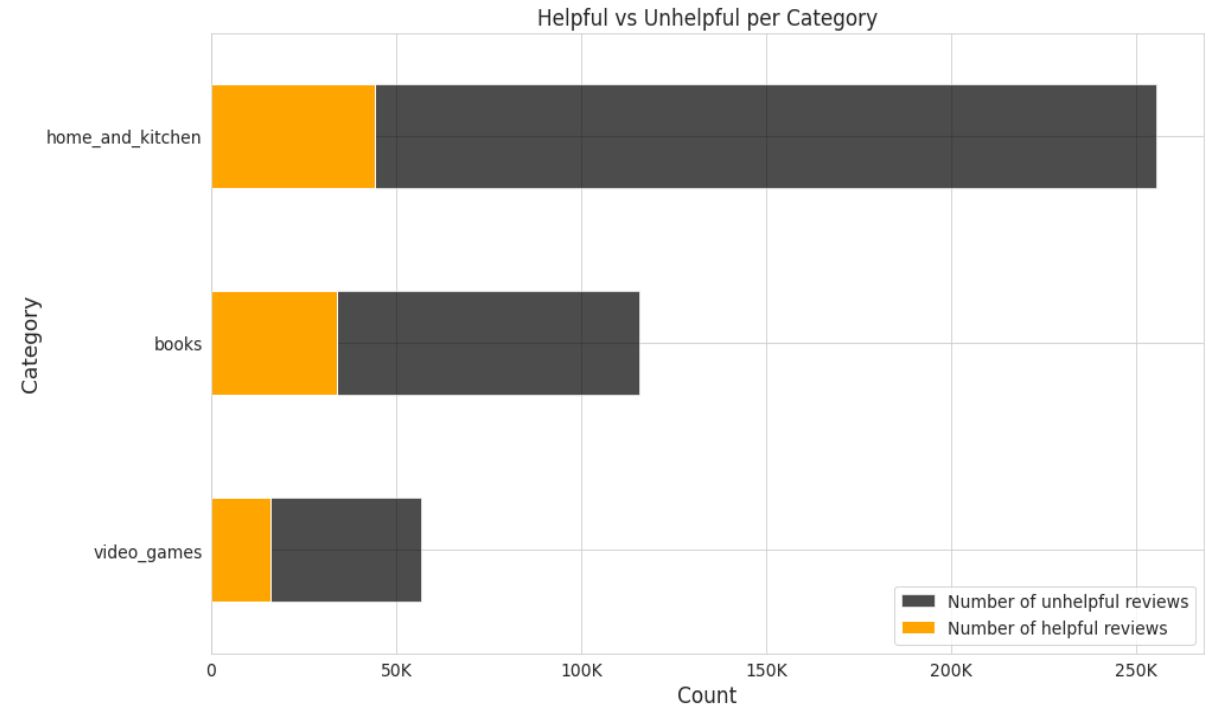
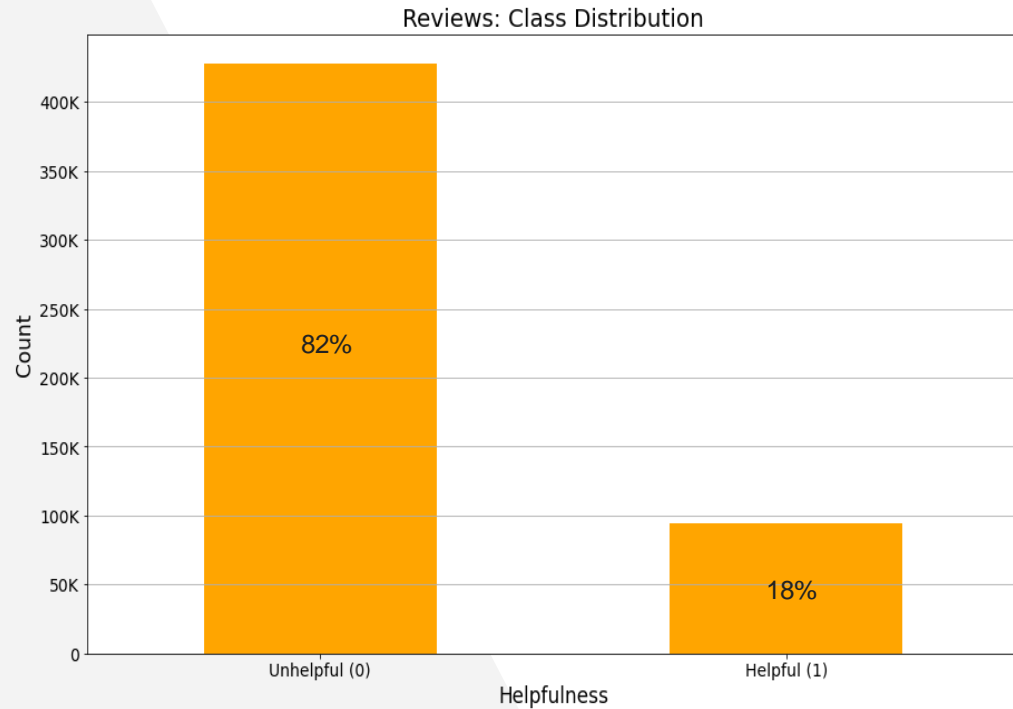
[https://www.brightlocal.com/research/local-consumer-review-survey/?SSAID=314743&SSCID=a1k5\\_3jwyh](https://www.brightlocal.com/research/local-consumer-review-survey/?SSAID=314743&SSCID=a1k5_3jwyh)



# 02

## Exploratory Data Analysis

# EDA – Target Class Split



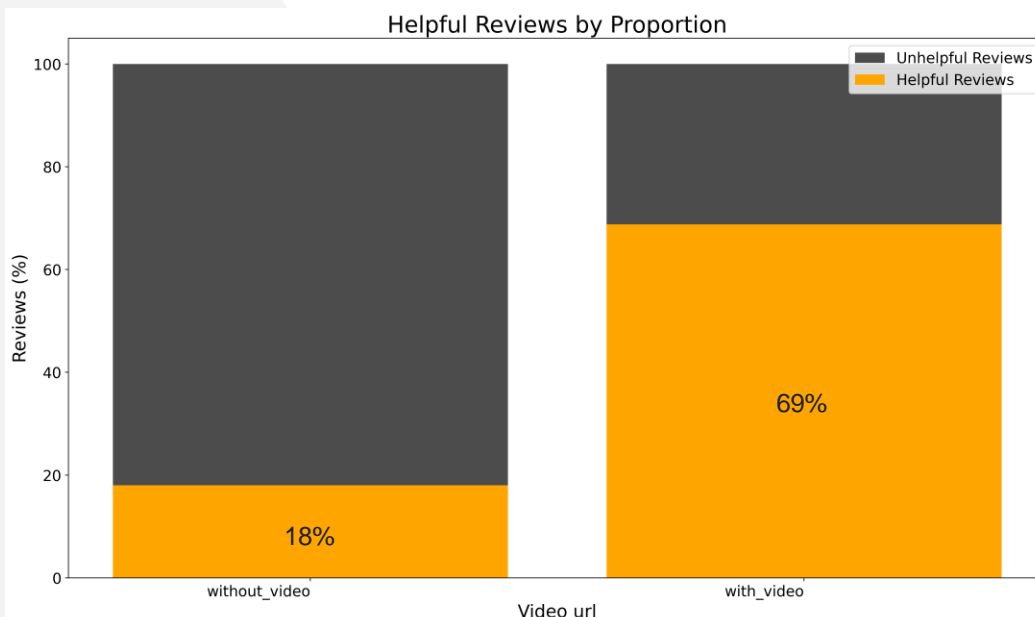
- ▶ Highly imbalanced - 18% of total reviews were helpful
- ▶ By category, 14%, 22.7%, and 22.0% from home and kitchen, books and videogames reviews were voted helpful by users respectively

# EDA –Video Reviews

Video Url Flag	Count	%
1	1,876	0.05%
0	3,485,455	99.95%

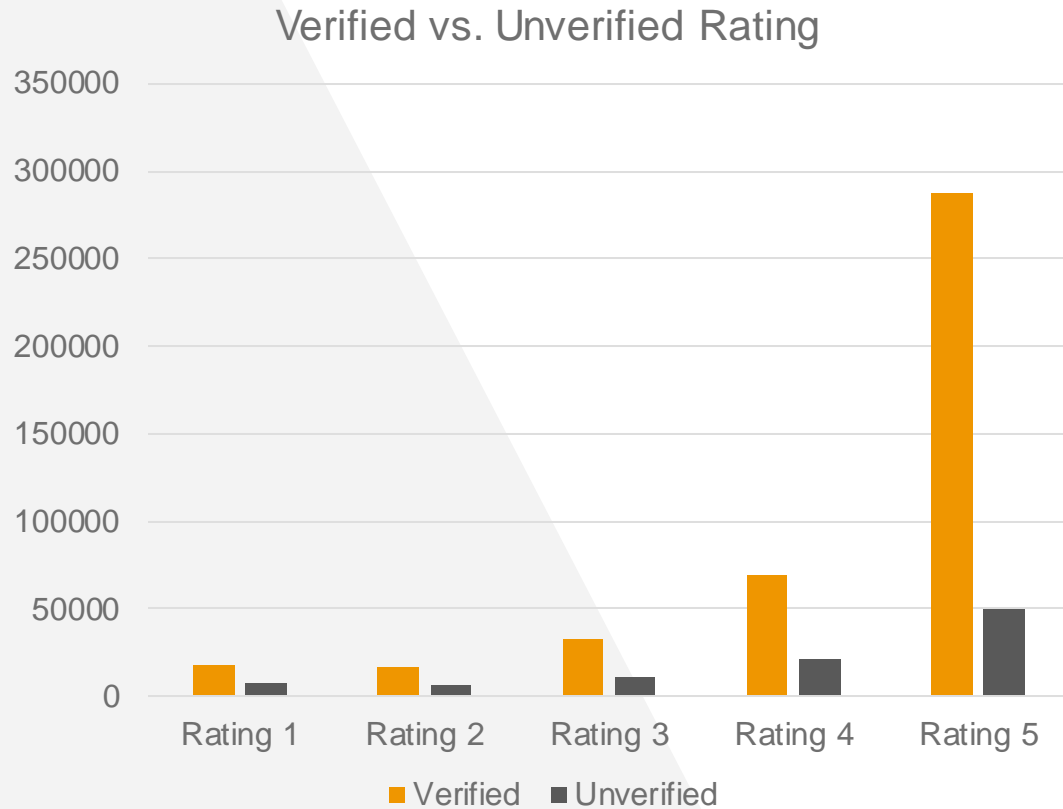
## Insights:

- ▶ About 0.05% of reviews contained video links
- ▶ About 70% of the reviews with videos were helpful
- ▶ Videos and actual images of products from unbiased users can help other users to make purchase decisions



reviewerID	review_summary
A18961KK4R3Y6A	<div id="video-block-R06P1T4N05ILD" class="a-section a-spacing-small a-spacing-top-mini video-block"></div><input type="hidden" name="" value="https://images-na.ssl-images-amazon.com/images/I/E1ZmFy4IRS.mp4" class="video-url"><input type="hidden" name="" value="https://images-na.ssl-images-amazon.com/images/I/11zL7xqwCaS.png" class="video-slate-img-url">&nbsp;Watch it in higher quality here: <a href="https://www.youtube.com/watch?v=qHood4BRRLw">https://www.youtube.com/watch?v=qHood4BRRLw</a> I am giving away a brand new copy of the collector's guide.
A2300Z6DNS1SWC	<div id="video-" type="hidden" <input type="h slate-img-url"> like to upload fo
A2W5AR8QG9VXV9	<div id="video-" type="hidden" <input type="h slate-img-url"> be reviewing ot
A2300Z6DNS1SWC	<div id="video-" type="hidden" <input type="h slate-img-url"> is rated mature

# EDA – Unverified Reviews

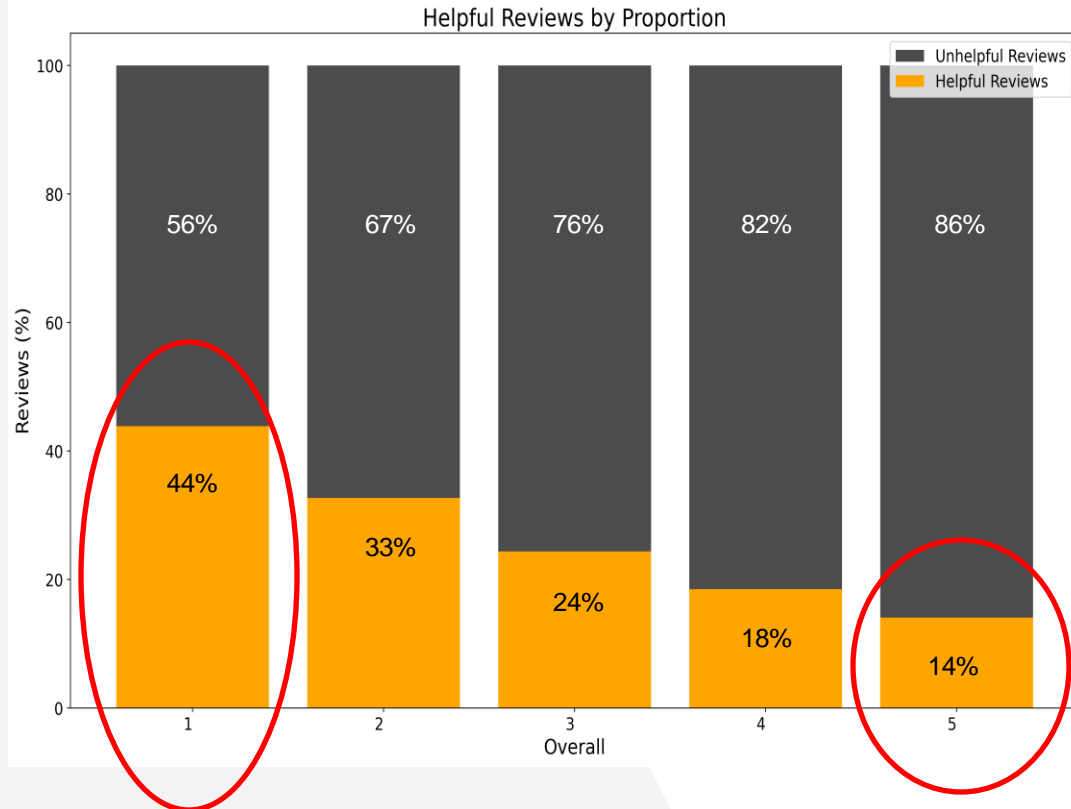


## Insights:

- ▶ Most accounts tend to give 5 stars to ratings
- ▶ Unverified accounts tend to review 5 stars the most: participative reviewers or bots?
- ▶ Sellers need to double think of their CRM strategy to deal with non-buyer reviewers; amazon needs to handle review bots at the same time



# EDA – Helpful reviews per rating

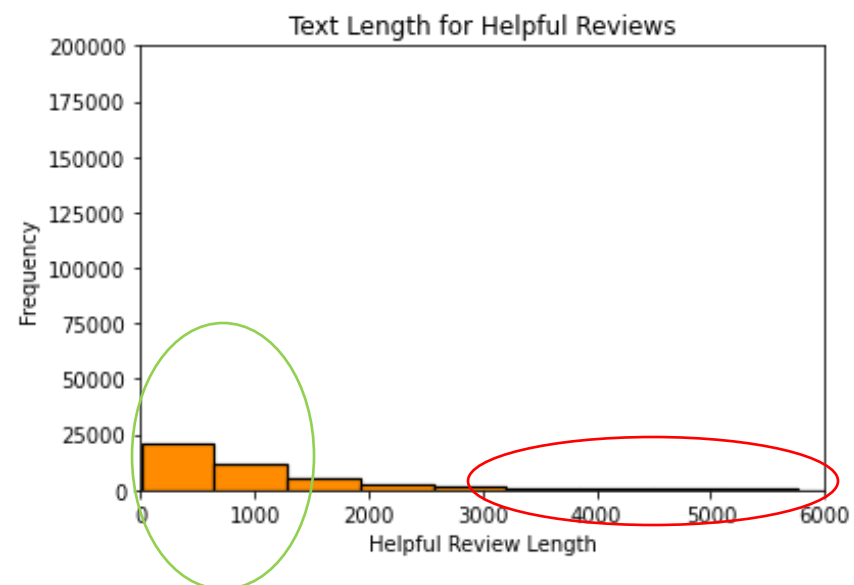
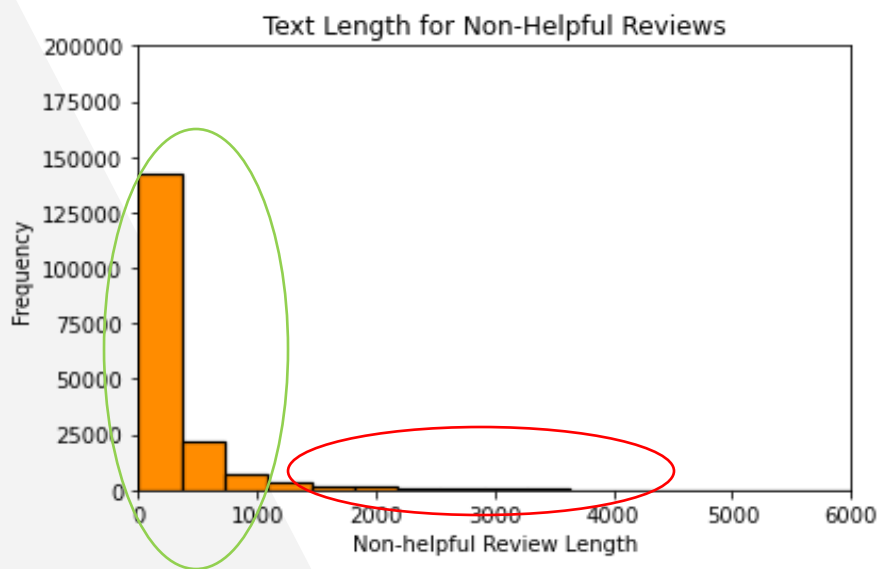


- ▶ On the scale 1 to 5, the lower the rating, the higher the % of helpful reviews
- ▶ Potential skewness as users tend to write more about their negative experience when they are dissatisfied
- ▶ Higher ratings tend to have the least % of helpful reviews as users tend to spend less time writing reviews when their experience is positive





# EDA-Text Length of Helpful Reviews

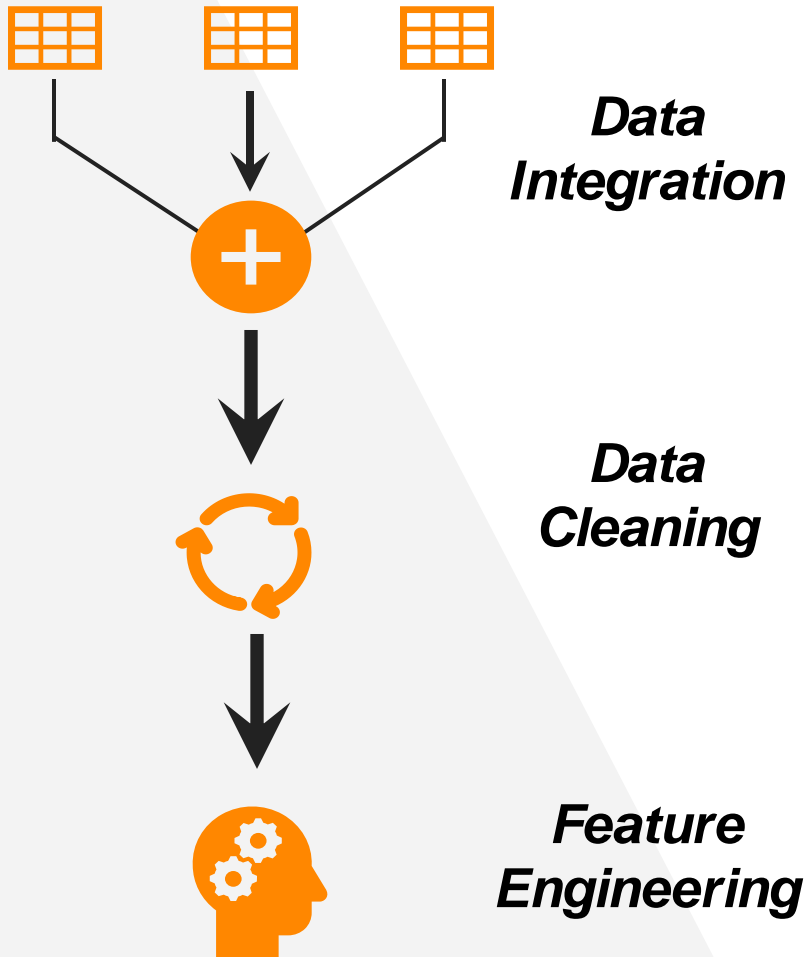


- ▶ On average, helpful reviews have higher text length both for review & summary
- ▶ Reviews that are not voted as 'helpful' are concentrated within the 1,000 characters in terms of text length

Label	Average Text Length
1	1,047
0	298

# 03 Modelling Journey

# Data Cleaning and Feature Engineering



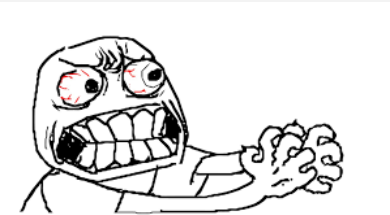
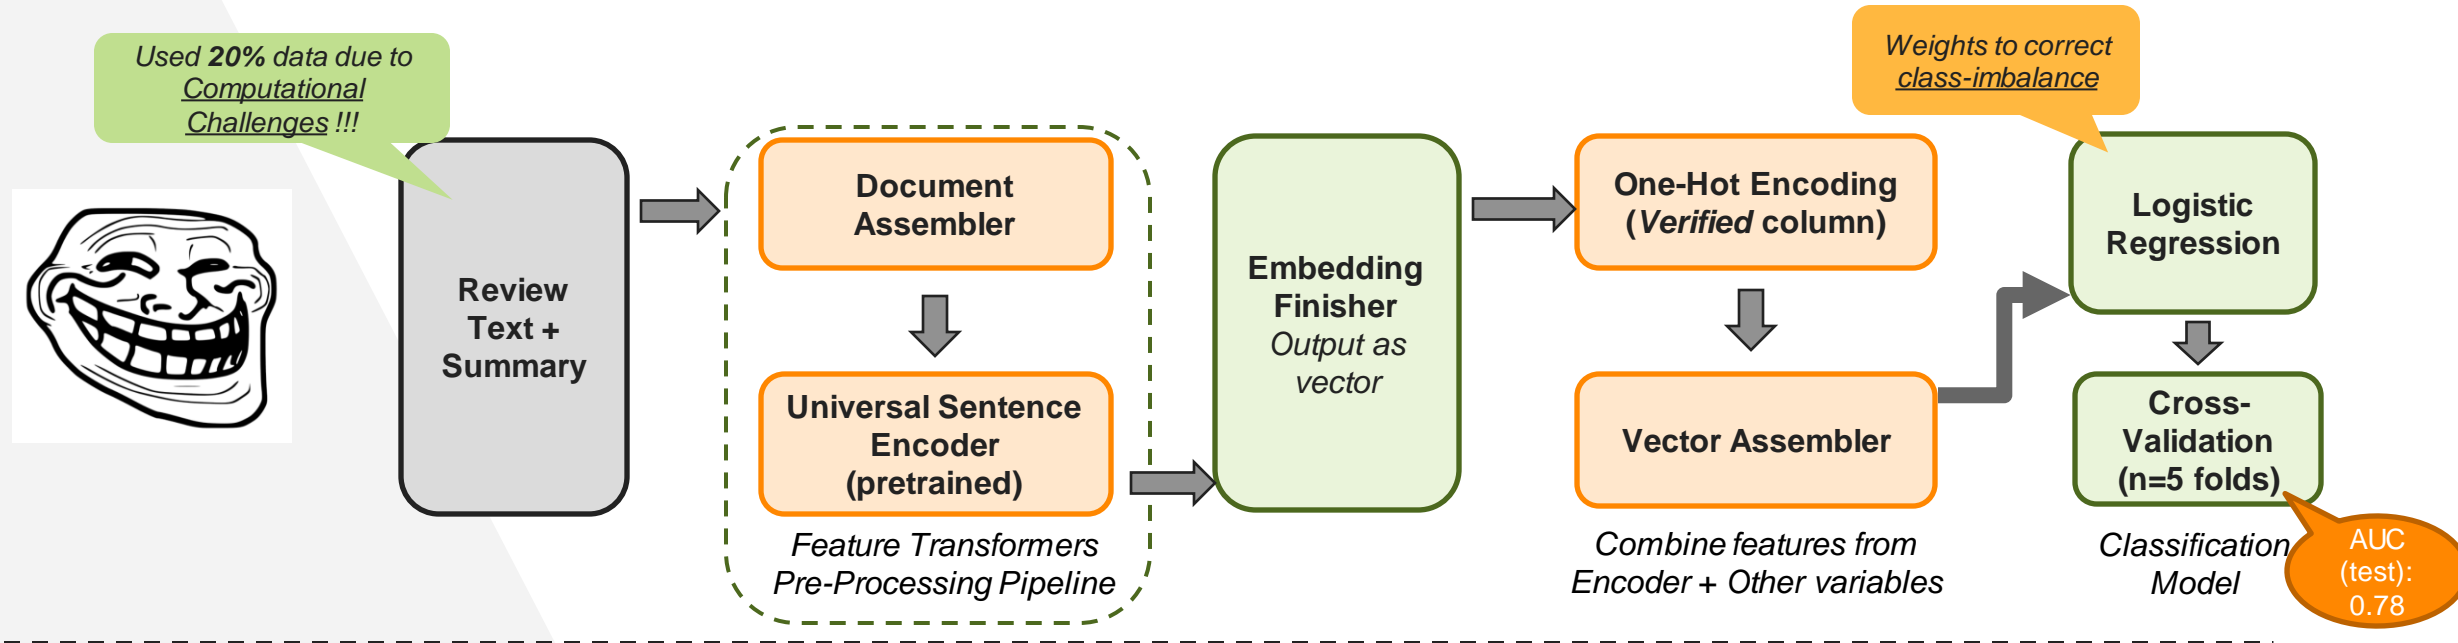
- Sourced datafiles from the server
- Merged three files

- Removed rows with < 10 characters in “reviewText”
- Merged text from “reviewText” and “summary”
- Removed special characters, weblink and numbers
- Converted text to lowercase
- Removed non-English instances

*Introduced following new features:*

- # of characters in review text
- # of words in review text
- Flag for review with more than 1000 characters
- Flag for reviews with video and image URLs

# Machine Learning Pipeline



## Transformers

- BERT
- XL Net
- ELMo

## Annotators

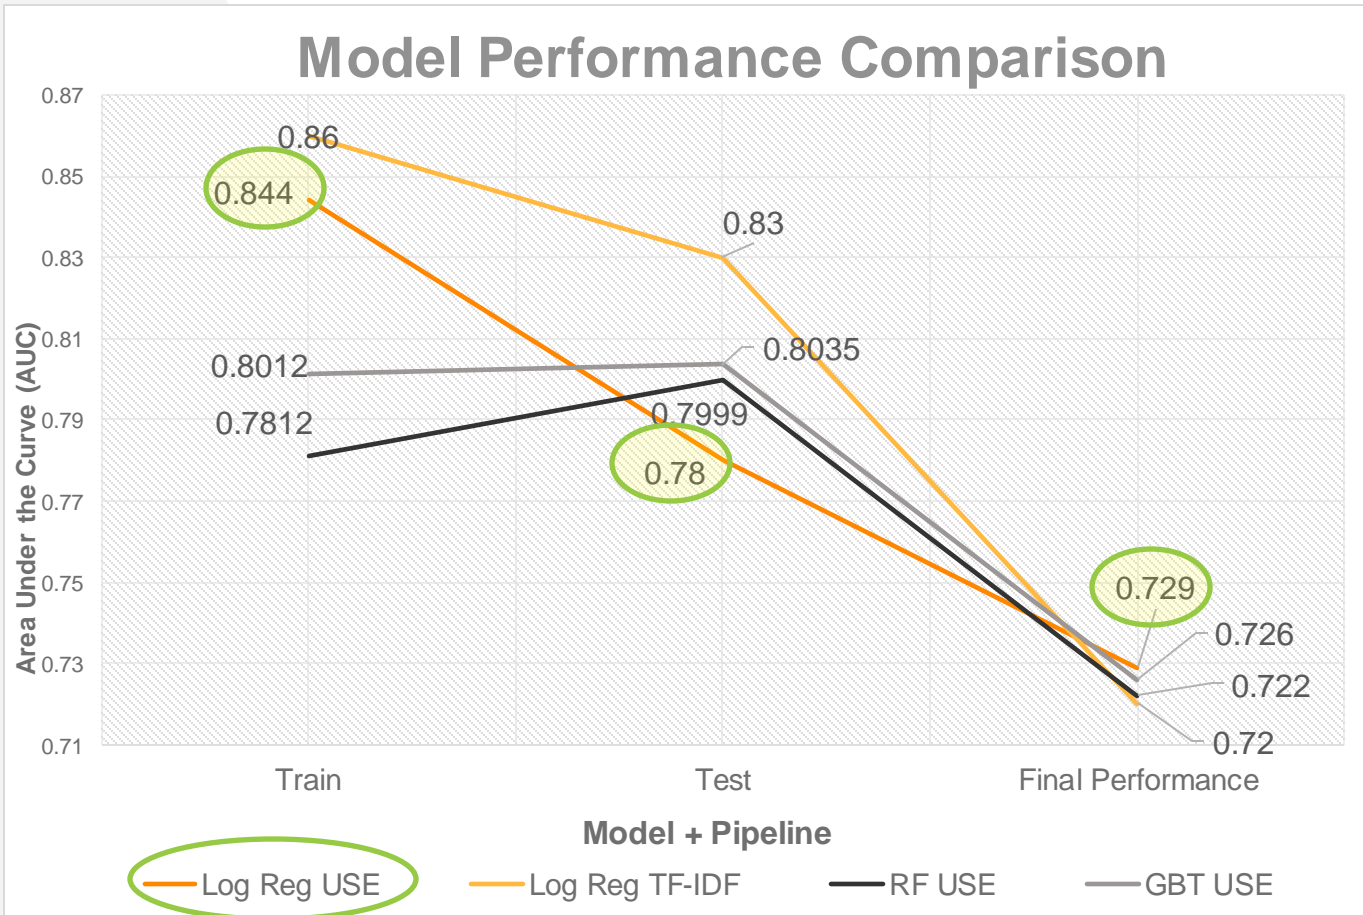
- ClassifierDL
- NorvigSweating Spellchecker
- Sentiment Detector

## Estimator

- RandomForest
- Gradient Boost
- Factorization Machines

# Modelling Process – Model + Pipeline Results Comparisons

## Model + Pipeline Comparison



### Top Performing Model

- Logistic Regression with Universal Sentence Encoder Pipeline

### Key Observations

- Significant overfitting of all models

# Modelling Process – Best Model Performance

## Classification Report

Class	Precision	Recall	F1-Score	Support
Non-Helpful (0)	0.87	0.96	0.91	57,152
Helpful (1)	0.64	0.34	0.45	12,579
-----	-----	-----	-----	-----
Accuracy			0.85	69,731
Macro Avg	0.75	0.65	0.68	69,731
Weighted Avg (Micro)	0.83	0.85	0.83	69,731

## Confusion Matrix

Truth Class	Predicted Class	
	0	1
0	54,726	2,426
1	8,265	4,314

## Current Kaggle Score/Position

- AUC: 0.7299
- Position: 13

# 04

## Project Cost Estimate

# Project Cost Estimate

## Compute Costs

- Price/Hour = \$1.173
- Hours used = 300
- Total Cost
- = **\$351.90**

## People Costs

- Data Scientist:  
\$120k | 20hrs
- Data Engineer:  
\$100k | 30hrs
- Analyst:  
\$80k | 80hrs
- Total Wages = **\$5,673**

## Infrastructure Costs

- Azure Data Lake  
Storage
- 5GB Data
- \$0.2304/GB
- Total Cost = **\$1.15**



**Grand Total Cost: \$6,026.13**



# 05

## Conclusion & Next Steps

# Conclusion & Next Steps

## In Summary

- ▶ Helpful reviews often come from unhappy customers
- ▶ Length of reviews matter!
- ▶ Best Model: Logistic Regression with USE



## What's Next?

- ▶ Main goal: Improve customer experience
- ▶ How? A/B test model to determine effectiveness

# Appendix

## ▶ Feature Engineering

```
#FE
import pyspark.sql.functions as f
#Adding in length of reviewText & Summary as features

df = df.withColumn('rt_length', f.length('reviewText'))
df = df.withColumn('sum_length', f.length('summary'))

df.show(10)
```

# Appendix

## ► Stratified Sampling

```
#Now, lets scale this down to 10% of our total data to train our algorithms
#Take a stratified sample

df = df.sampleBy('label', fractions={0:0.1, 1:0.1}, seed=47)

display(df.groupBy('label').count())
```

(2) Spark Jobs

df: pyspark.sql.dataframe.DataFrame = [reviewID: long, overall: double ... 13 more fields]

	label	count
1	1	62746
2	0	286301

# Appendix

## ▶ Feature Engineering

```
#Adding in rt_length indicator

df = df.withColumn('rt_length_ind', f.when(f.col('rt_length') >= 1000, 1).otherwise(0))

df.show(10)

#df.withColumn('D', f.when(f.col('B') > 0, "Yes").otherwise("No")).show()
```

# Appendix

## ► Feature Engineering

```
#Count total number of sentences
def countWordsInEachSentences(array):
    return [len(x.split()) for x in array]

countWordsSentences = f.udf(lambda x: countWordsInEachSentences(x.split('. ')))

df = df.withColumn("word_count", countWordsSentences(df['reviewText']))

df.show(10)
```

# Appendix

## ► Language Detection Pipeline

```
#Checking reviewText for non english before splitting
import sparknlp
from sparknlp.annotator import *
from sparknlp.common import *
from sparknlp.base import *
from sparknlp.pretrained import PretrainedPipeline

model_name = 'ld_wiki_tatoeba_cnn_375'

documentAssembler = DocumentAssembler() \
    .setInputCol("reviewText") \
    .setOutputCol("document")

sentence_detector = SentenceDetector() \
    .setInputCols(["document"]) \
    .setOutputCol("sentence")

languageDetector = LanguageDetectorDL.pretrained(model_name) \
    .setInputCols("sentence") \
    .setOutputCol("language") \
    .setThreshold(0.5) \
    .setCoalesceSentences(True)

nlpPipeline = Pipeline(stages=[ documentAssembler,
                                sentence_detector,
                                languageDetector
                                ])

langpipelineModel = nlpPipeline.fit(df)
```

# Appendix

## ► Universal Sentence Encoder Pipeline 1/2

```
#Downloading & Initializing the Universal Sentence Encoder + Importing relevant libraries
import sparknlp
from sparknlp.annotator import *
from sparknlp.common import *
from sparknlp.base import *
from pyspark.ml import Pipeline
from pyspark.ml.feature import CountVectorizer, HashingTF, IDF, OneHotEncoder, StringIndexer, VectorAssembler, SQLTransformer
from pyspark.sql.functions import udf, explode
import pandas as pd

useEmbeddings = UniversalSentenceEncoder.pretrained()\
    .setInputCols("document")\
    .setOutputCol("use_embeddings")
```



# Appendix

## ► Universal Sentence Encoder Pipeline 2/2

```
#Creating the pipeline
document_assembler = DocumentAssembler() \
    .setInputCol("reviewText") \
    .setOutputCol("document")

loaded_useEmbeddings = UniversalSentenceEncoder.load('/root/cache_pretrained/tfhub_use_en_2.4.0_2.4_1587136330099') \
    .setInputCols("document") \
    .setOutputCol("use_embeddings")

embeddings_finisher = EmbeddingsFinisher() \
    .setInputCols(["use_embeddings"]) \
    .setOutputCols(["finished_use_embeddings"]) \
    .setOutputAsVector(True) \
    .setCleanAnnotations(False)
```

# Appendix

## ▶ Extracting Universal Sentence Encoder Embeddings

```
#Viewing embedding output
use_df_train.select('finished_use_embeddings').show(3)

(1) Spark Jobs
-----+
finished_use_embeddings|
-----+
[[-0.056037165224...|
[[-0.005085167475...|
[[-0.045702368021...|
-----+
y showing top 3 rows

Command took 16.38 seconds -- by 20mg46@queensu.ca at 9/30/2021, 4:01:21 PM on MMA2022W-Kipling
6

#Adding in a features column based on embeddings
use_df_train= use_df_train.withColumn("fue", explode(use_df_train.finished_use_embeddings))

#use_df_train = use_df_train.withColumn('xln', explode(use_df_train.finished_xlnet_embeddings))
```

# Appendix

## ► Pipeline 2

```
#Now create new pipeline to encode verified and pass fue, overall and verified together into one column called "features"
#Now OHE the Verified feature and add in the overall feature as well

verified_indexer = StringIndexer(inputCol=('verified'), outputCol=('verified_indexer'))

ohe_verified_indexer = OneHotEncoder(inputCol=('verified_indexer'), outputCol=('verified_indexer_vec'))

#overall_indexer = StringIndexer(inputCol=('overall'), outputCol=('verified_overall'))

#ohe = OneHotEncoder()\
# .setInputCols(['verified'])\
# .setOutputCols(['verified_0'])

#assembler = VectorAssembler(inputCols=['fue', 'verified_indexer_vec', 'overall', 'rt_length', 'sum_length', 'rt_length_ind'], outputCol='features')
assembler = VectorAssembler(inputCols=['fue', 'verified_indexer_vec', 'overall', 'rt_length', 'rt_length_ind'], outputCol='features')

pre_proc_pipeline = Pipeline(stages=[verified_indexer, ohe_verified_indexer, assembler])

pre_proc_model = pre_proc_pipeline.fit(use_df_train)
```

# Appendix

## ► Logistic Regression Model

```
#Log Reg TV
import mlflow
import mlflow.spark
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, auc
import pandas as pd
from pyspark.ml.classification import LogisticRegression, GBTClassifier
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.tuning import CrossValidator, TrainValidationSplit, ParamGridBuilder

#Define Model
lr = LogisticRegression(labelCol='label', featuresCol='features')

#Define Evaluator
evaluator = BinaryClassificationEvaluator(labelCol='label', metricName='areaUnderROC')

lr_paramGrid = ParamGridBuilder() \
    .addGrid(lr.regParam, [0.0, 0.5, 0.7]) \
    .addGrid(lr.maxIter, [10,25,50,100]) \
    .build()

lr_tvs = TrainValidationSplit(estimator=lr, estimatorParamMaps=lr_paramGrid, evaluator=evaluator,
                             parallelism=4, seed=42)
lr_grid_model = lr_tvs.fit(train_preprocessed)
```

# Appendix

## ► Logistic Regression Model 2

```
#Logistic Regression
import mlflow
import mlflow.spark
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, auc
import pandas as pd
from pyspark.ml.classification import LogisticRegression, GBTCClassifier
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.tuning import CrossValidator, TrainValidationSplit, ParamGridBuilder

#Let's start off with a Logistic Regression model
lr = LogisticRegression(elasticNetParam=0.0)

#Define our evaluator
evaluator = BinaryClassificationEvaluator(labelCol='label', metricName='areaUnderROC')

#Define a hyperparameter search space
lr_paramGrid = ParamGridBuilder() \
    .addGrid(lr.regParam, [0.0, 0.5, 0.7]) \
    .addGrid(lr.maxIter, [10,25,50,100]) \
    .build()

lr_grid = CrossValidator(estimator=lr,
                        estimatorParamMaps=lr_paramGrid,
                        evaluator=evaluator,
                        numFolds=5)

# Run cross-validation, and choose the best set of parameters.
with mlflow.start_run():
    lr_grid_model = lr_grid.fit(train_preprocessed) #Running cv on the training dataset; will return the best model it found

    #Evaluate the best model's performance on the test set and log the results
    #test_metric = evaluator.evaluate(lr_grid_model.transform(use_df_test))
    #mlflow.log_metric('test_' + evaluator.getMetricName(), test_metric)

    #Log the best model
    mlflow.spark.log_model(spark_model=lr_grid_model.bestModel, artifact_path='best-model')
```

# Appendix

## ▶ Random Forest Model

```
#RF TV
import mlflow
import mlflow.spark
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, auc
import pandas as pd
from pyspark.ml.classification import LogisticRegression, GBClassifier, RandomForestClassifier
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.tuning import CrossValidator, TrainValidationSplit, ParamGridBuilder

#Define Model
rf = RandomForestClassifier(labelCol='label', featuresCol='features')

#Define Evaluator
evaluator = BinaryClassificationEvaluator(labelCol='label', metricName='areaUnderROC')

rf_paramGrid = ParamGridBuilder() \
    .addGrid(rf.maxBins, [10,20,30]) \
    .addGrid(rf.maxDepth, [4,6,8]) \
    .addGrid(rf.impurity, ['gini', 'entropy']) \
    .build()

rf_tvs = TrainValidationSplit(estimator=rf, estimatorParamMaps=rf_paramGrid, evaluator=evaluator,
                             parallelism=4, seed=42)
rf_grid_model = rf_tvs.fit(train_preprocessed)
```

# Appendix

## Gradient Boost Tree Classifier

```
#Gradient Boosted Tree Classifier
import mlflow
import mlflow.spark
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, auc
import pandas as pd
from pyspark.ml.classification import LogisticRegression, GBTClassifier
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.tuning import CrossValidator, TrainValidationSplit, ParamGridBuilder

#Define Model
gbt = GBTClassifier(labelCol='label', featuresCol='features')

#Define Evaluator
evaluator = BinaryClassificationEvaluator(labelCol='label', metricName='areaUnderROC')

gbt_paramGrid = ParamGridBuilder() \
    .addGrid(gbt.maxDepth, [3,5,10]) \
    .addGrid(gbt.maxBins, [2,12,24,32]) \
    .build()

gbt_tvs = TrainValidationSplit(estimator=gbt, estimatorParamMaps=gbt_paramGrid, evaluator=evaluator,
                              parallelism=4, seed=42)
gbt_model = gbt_tvs.fit(train_preprocessed)
```

# Appendix

## ► Compute Cost Calculations

Azure VM								
RAM	vCPUs	DBU Count	VM Price/CPU	VM Price	DBU Price	Total Price/Hour	Total Hours Used	Total Cost
28GB	4	1	0.1173	\$ 0.4690	\$ 0.7040	\$ 1.173	300	\$ 351.90



## Appendix

### ▶ Employee Wage Cost Calculations

Employees				
Title	Pay	Hours Used	Hourly Salary	Total Cost
Data Scientist	\$ 120,000	20	\$ 57.69	\$ 1,153.85
Data Engineer	\$ 100,000	30	\$ 48.08	\$ 1,442.31
Analyst	\$ 80,000	80	\$ 38.46	\$ 3,076.92
				\$ 5,673.08

# Appendix

## ► Infrastructure Cost Calculations

Infrastructure		
GBs Used	Price	Total
5	\$ 0.2304	\$ 1.15